

Generalization and Distributed Learning of GFlowNets

Tiago da Silva, Amauri Souza, Omar Rivasplata, Vikas Garg, Samuel Kaski, Diego Mesquita



KEYWORDS — GFlowNets, Distributed learning, PAC-Bayes

TL;DR

- we introduce the first non-vacuous generalization bounds for GFlowNets,
- we develop the first Azuma-type PAC-Bayesian bounds for understanding the generalization of GFlowNets under the light of Martingale theory,
- we demonstrate the harmful effect of the trajectory length on the proven learnability of a generalizable policy for GFlowNets,
- we introduce the first distributed algorithm for learning GFlowNets, Subgraph Asynchronous Learning, and show that it drastically accelerates learning convergence and mode discovery when compared against a centralized approach for relevant benchmark tasks

I. BACKGROUND: GFLOWNETS

GFlowNets are amortized algorithms for sampling from distributions over discrete and compositional objects (such as graphs).




Figure 1: A GFlowNet learns a **forward policy** on a state graph.

A **flow network** is defined over an extension \mathcal{S} of \mathcal{G} , which then represents the sink nodes. To navigate through this network and sample from \mathcal{G} in proportion to a **reward function** $R: \mathcal{G} \rightarrow \mathbb{R}_+$, a forward (resp. backward) policy $p_F(\tau)$ ($p_B(\tau|x)$) is used.

$$p_F(\tau) = \prod_{(s,s') \in \tau} p_F(s' | s) \text{ and } \sum_{\tau \sim g} p_F(\tau) = R(g). \quad (1)$$

To achieve this, we parameterize $p_F(\tau)$ as a neural network trained by minimizing

$$\mathcal{L}_{TB}(p_F) = \mathbb{E} \left[\left(\log \frac{p_F(\tau)Z}{p_B(\tau|x)R(x)} \right)^2 \right]. \quad (2)$$

for a given $p_B(\tau|x)$. GFlowNets can be trained in an **off-policy** fashion and the above expectation can be under any full-support distribution over trajectories.

II. BACKGROUND: PROBABLY APPROXIMATELY CORRECT BAYESIAN BOUNDS

Let \mathcal{L} be a loss function on a parameter space Θ , e.g., the squared loss. Also, let $\hat{\mathcal{L}}(\theta, \mathbf{X})$ be its empirical counterpart evaluated on a dataset \mathbf{X} .

PAC-Bayesian bounds. Given “prior” Q (independent of \mathbf{X}) and posterior P distributions over Θ , a PAC-Bayesian bound establishes an upper limit for the expectation of (unobserved) \mathcal{L} based on the (observed) $\hat{\mathcal{L}}$ and a complexity term φ and a confidence level δ ,

$$\mathbb{E}_{\theta \sim P}[\mathcal{L}(\theta)] \leq \mathbb{E}_{\theta \sim P}[\hat{\mathcal{L}}(\theta, \mathbf{X})] + \varphi(\delta, P, Q, |\mathbf{X}|). \quad (3)$$

When $\mathcal{L}(\theta) \leq B$ a.e., we refer to a bound as *vacuous* if

$$\mathbb{E}_{\theta \sim P}[\hat{\mathcal{L}}(\theta, \mathbf{X})] + \varphi(\delta, P, Q, |\mathbf{X}|) \geq B. \quad (4)$$

Otherwise, the bound is *non-vacuous*. Historically, the search for non-vacuous PAC-Bayesian bounds has been associated to the search for provably generalizable learning algorithms. In this regard, recent works have built upon the basic PAC-Bayesian inequalities to obtain theoretical guarantees for GANs, transformers, armed bandits, and variational autoencoders.

Data-dependent priors for PAC-Bayesian bounds. Often, φ involves the KL divergence between P and Q , which dominates the upper bound and commonly results in vacuously true statements. To circumvent this issue, we separate $\mathbf{X} = \mathbf{X}_\alpha \cup \mathbf{X}_{1-\alpha}$ into disjoint and independent subsets. A posterior P is learned on $\mathbf{X}_{1-\alpha}$ through conventional methods and a prior Q is subsequently learned on \mathbf{X}_α by minimizing the PAC-Bayesian upper bound,

$$Q^* = \operatorname{argmin}_Q \mathbb{E}_{\theta \sim P}[\hat{\mathcal{L}}(\theta, \mathbf{X}_\alpha)] + \varphi(\delta, P, Q, \alpha |\mathbf{X}|). \quad (5)$$

III. NON-VACUOUS GENERALIZATION BOUNDS FOR GFLOWNETS

There are four ingredients for a PAC-Bayesian bound: a bounded risk functional \mathcal{L} , a prior distribution Q , a posterior distribution P , and a learning algorithm. In alignment with the broader literature, we use a diagonal Gaussian distribution for both P and Q with fixed (small) variance. For learning, we use SGD.

A bounded risk functional for GFlowNets. In a recent work, we demonstrated that Flow Consistency in Subgraphs (FCS) is a sound and tractable learning objective for GFlowNets.

$$\text{FCS}(\pi, p_\top) = \mathbb{E}_{\mathcal{B}}[\text{TV}(\pi^{\mathcal{B}}, p_\top^{\mathcal{B}})], \quad (6)$$

in which $\pi \propto R$ is the target distribution and

$$p_\top(x) = \sum_{\tau \text{ finishing at } x} p_F(\tau) \quad (7)$$

is the probability of $x \in \mathcal{G}$ under p_F ; the expectation is under a distribution of random independent subsets of \mathcal{G} . Intuitively, FCS measures the total variation TV between π and p_\top on random subgraphs of the underlying flow network.

The stochastic and bounded nature of FCS make it a suitable candidate for pursuing a PAC-Bayesian analysis of GFlowNets. We will refer to FCS by L and to its empirical counterpart by \hat{L} to emphasize its use as a risk functional for assessing the generalization of GFlowNets.

Non-vacuous generalization bounds. Let \mathcal{T}_n be a n -sized set of trajectories sampled from a stationary distribution. Also, let P and Q be distributions over Θ . Then,

$$L_{\text{FCS}}(P) \leq \hat{L}_{\text{FCS}}(P, \mathcal{T}_{1-\alpha}) + \min \left(\left\{ \eta + \sqrt{\eta(\eta + 2\hat{L}_{\text{FCS}}(P, \mathcal{T}_{1-\alpha}))} \right\}, \sqrt{\frac{\eta}{2}} \right) \quad (8)$$

in which the complexity term η is, for chosen $\alpha \in (0, 1)$,

$$\eta := \frac{\text{KL}(P \parallel Q) + \log \frac{2\sqrt{(1-\alpha)n}}{\delta}}{[(1-\alpha)n]}. \quad (9)$$

We optimize Equation 8 to obtain data-dependent priors Q over Θ . Figure 2 shows the resulting bounds are non-vacuous. These are the first positive and rigorous results regarding the generalization GFlowNets in the literature.




Figure 2: Non-vacuous generalization bounds for GFlowNets.

Oracle generalization bounds for GFlowNets. Let \mathcal{L} be the within-trajectory detailed balance loss function and assume that $\mathcal{L} \leq U$ a.e.. Additionally, define t_m as the maximum trajectory length within the flow network and T as a budget for the number of transitions.

$$\mathbb{E}_{\theta \sim P}[\mathcal{L}(\theta)] \leq \frac{1}{\beta} \mathbb{E}_{\theta \sim P}[\hat{\mathcal{L}}(\theta)] + \alpha_{T,n} \left(\text{KL}(P \parallel Q) + \log \frac{2}{\delta} \right) + \frac{\log t_m}{\beta T \lambda} + \gamma \frac{\lambda 2U^2}{\beta T} \quad (10)$$

in which $\beta \in (0, 1)$, $\lambda > 0$, and

$$\alpha_{T,n} = \frac{U}{2\beta(1-\beta)n} + \frac{1}{\beta T \lambda}. \quad (11)$$

IV. SUBGRAPH ASYNCHRONOUS LEARNING (SAL)

Equation 10 demonstrates that the larger trajectory length t_m play a key role in constraining the generalization potential of GFlowNets. To mitigate this effect, we propose a distributed divide-and-conquer learning algorithm that breaks up the state graph into smaller subgraphs and learns a GFlowNet for each subgraph. The resulting GFlowNets are aggregated in a final, efficient step. We refer to this approach as **Subgraph Asynchronous Learning (SAL)**.

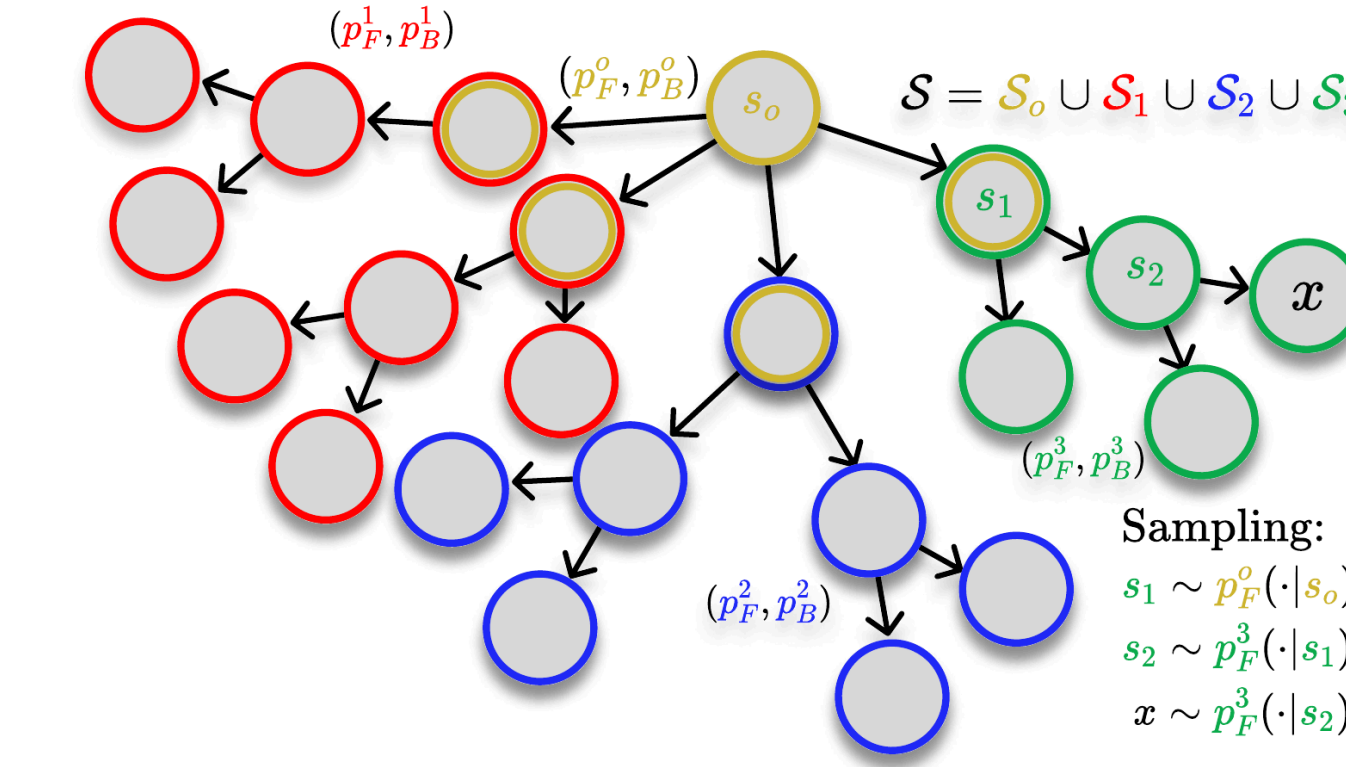


Figure 3: An illustration of SAL.

Algorithm 1 Subgraph Asynchronous Learning

- 1: $\mathcal{S} = \mathcal{S}_o \cup \bigcup_{j=1}^m \mathcal{S}_j \triangleright$ Fixed-horizon partition
- 2: $\mathcal{I}_j = \mathcal{S}_j \cap \mathcal{S}_o$ for $j \in \{1, \dots, m\}$
- 3: \triangleright Local training
- 4: **for** $j \in \{1, \dots, m\}$ **do**
- 5: \triangleright Minimize $\mathcal{L}_{\text{ATB}}^j$ in \mathcal{S}_j with SGD
- 6: $(p_F^j, F_j) = \arg \min_{p_F, F} \mathcal{L}_{\text{ATB}}^j(p_F, F)$
- 7: **end for**
- 8: $R^o: x \mapsto \mathbb{1}_{\{x \in \mathcal{X}\}} R(x) + \sum_{j=1}^m \mathbb{1}_{\{x \in \mathcal{I}_j\}} F_j(x)$
- 9: $(p_F^o, F_o) = \arg \min_{p_F, F} \mathcal{L}_{\text{TB}}(p_F, F, R^o)$
- 10: **return** $\{(p_F^o, F_o)\} \cup \bigcup_{1 \leq j \leq m} \{(p_F^j, F_j)\}$

SAL. Let $\{(p_F^1, F_1), \dots, (p_F^m, F_m)\}$ be m GFlowNets defined over each of the m components of the partition defining SAL. Also, let q_j be a distribution over the initial states and p_E^j be an distribution over trajcotires within the j th component \mathcal{S}_j . Then, each (p_F^j, F_j) is learned by minimizing the *amortized trajectory balance loss* over \mathcal{S}_j

$$\mathcal{L}_{\text{ATB}}^j(p_F^j, F_j) = \mathbb{E}_{s \sim q_j} \mathbb{E}_{\tau \sim p_E^j(\cdot | s)} \left[\left(\log \frac{F_j(s)p_F^j(\tau | s)}{R(x)p_B^j(\tau | x)} \right)^2 \right], \quad (12)$$

which replaces Z by the flow function F^j in the conventional \mathcal{L}_{TB} .

Our experimental results in Figure 5 and Figure 6 show that SAL often accelerate training convergence and increase the mode-finding capability of the GFlowNet. These observations are in alignment with our theoretical analysis in Equation 10 (Figure 5) and corroborate the intuition that a divide-and-conquer approach enhances the exploration of the learning agent (Figure 6).




Figure 5: SAL improves learning convergence for the hypergrid environment.




Figure 6: SAL drastically accelerates mode-finding for benchmark tasks.

Recursive SAL. SAL can be hierarchically extended to accommodate nested partitions of the state graph. This is illustrated in Figure [ref], and the resulting method is referred to as Recursive SAL. Notably, the depth of the nested partition characterizes a trade-off between the number of trainable models and the difficult of the problem that each model solves. The balance between these factors should be addressed in a case-by-case basis in future endeavors.

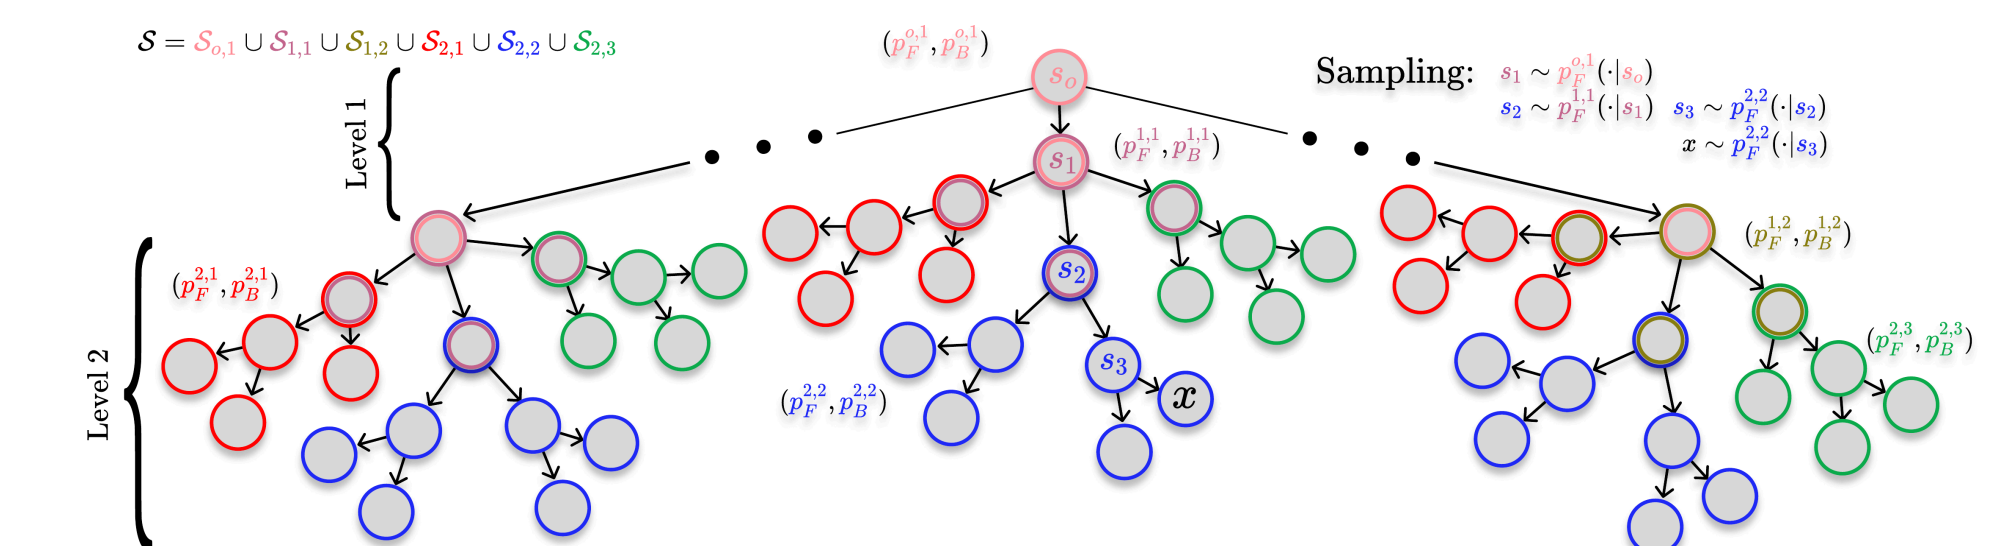


Figure 7: Illustration of Recursive SAL as an hierarchical extension of SAL.